Finding Scientific Gems with Google

P. Chen,[1, *] H. Xie,[2, 3, †] S. Maslov,[3, ‡] and S. Redner[1, §]

¹Center for Polymer Studies and and Department of Physics, Boston University, Boston, MA, 02215
²New Media Lab, The Graduate Center, CUNY, New York, NY, 10016
³Department of Condensed Matter Physics and Materials Science,
Brookhaven National Laboratory, Upton, NY, 11973

We apply the Google PageRank algorithm to assess the relative importance of all publications in the Physical Review family of journals from 1893–2003. While the Google number and the number of citations for each publication are positively correlated, outliers from this linear relation identify some exceptional papers or "gems" that are universally familiar to physicists.

PACS numbers: 02.50.Ey, 05.40.-a, 05.50.+q, 89.65.-s

I.   INTRODUCTION

With the availability of electronically available citation data, it is now possible to undertake comprehensive studies of citations that were unimaginable just a few years ago. In most previous studies of citation statistics, the metric used to quantify the importance of a paper is its number of citations. In terms of the underlying citation network, in which nodes represent publications and directed links represent a citation from a *citing* article to a *cited* article, the number of citations to an article translates to the in-degree of the corresponding node. The distribution of in-degree for various citation data sets has a broad tail [1] that is reasonably approximated by a power law [2, 3, 4].

While the number of citations is a natural measure of the impact of a publication, we probably all have encountered examples where citations do not seem to provide a full picture of the influence of a publication. We are thus motivated to study alternative metrics that might yield a truer measure of importance than citations alone. Such a metric already exists in the form of the Google PageRank algorithm [5]. A variant of the PageRank algorithm was recently applied to better calibrate the impact factor of scientific journals [6].

In this work, we apply Google PageRank to the Physical Review citation network with the goal of measuring the importance of individual scientific publications published in the APS journals. This network consists of 353,268 nodes that represent all articles published in the Physical Review family of journals from the start of publication in 1893 until June 2003, and 3,110,839 links that represent all citations *to* Physical Review articles *from* other Physical Review articles. As found previously [4], these internal citations represent 1/5 to 1/3 of all citations for highly-cited papers. This range provides a sense of the degree of completeness of the Physical Review ci-

tation network.

With the Google PageRank approach, we find a number of papers with a modest number of citations that stand out as exceptional according to the Google PageRank analysis. These exceptional publications, or gems, are familiar to almost all physicists because of the very influential contents of these articles. Thus the Google PageRank algorithm seems to provide a new and useful measure of scientific quality.

II.   THE GOOGLE PAGERANK ALGORITHM

To set the stage for our use of Google PageRank to find scientific gems, let us review the elements of the PageRank algorithm. Given a network of N nodes $i = 1, 2, \ldots, N$, with directed links that represent references from an initial (citing) node to a target (cited) node, the Google number G_i for the i^{th} node is defined by the recursion formula [5]:

$$G_i = (1 - d) \sum_{j \text{ nn } i} \frac{G_j}{k_j} + \frac{d}{N} . \qquad (1)$$

Here the sum is over the neighboring nodes j in which a link points to node i. The first term describes propagation of the probability distribution of a random walk in which a walk at node j propagates to node i with probability $1/k_j$, where k_j is the out-degree of node j. The second term describes the uniform injection of probability into the network in which each node receives a contribution d/N at each step.

Here d is a free parameter that controls the performance of the Google PageRank algorithm. The prefactor $(1-d)$ in the first term gives the fraction of random walks that continue to propagate along the links; a complementary fraction d is uniformly re-injected into the network, as embodied by the second term.

We suggest that the Google number G_i of paper i, defined by Eq. (1), is a better measure of importance than the number of citations alone in two aspects: i) being cited by influential papers contributes more to the Google number than being cited by unimportant papers; ii) being cited by a paper that itself has few references

gives a larger contribution to the Google number than being cited by a paper with hundreds of references. The Google number of a paper can be viewed as a measure of its influence that is then equally exported to all of its references. The parameter $d > 0$ prevents all of the influence from concentrating on the oldest papers.

In the original Google PageRank algorithm of Brin and Page [5], the parameter d was chosen to be 0.15. This value was prompted by the anecdotal observation that an individual surfing the web will typically follow of the order of 6 hyperlinks, corresponding to a leakage probability $d = 1/6 \simeq 0.15$, before becoming either bored or frustrated with this line of search and beginning a new search. In the context of citations, we conjecture that entries in the reference list of a typical paper are collected following somewhat shorter paths of average length 2, making the choice $d = 0.5$ more appropriate for a similar algorithm applied to the citation network. The empirical observation justifying this choice is that approximately 50% of the articles [9] in the reference list of a given paper A have at least one citation B $\to$ C to another article C that is also in the reference list of A (Fig. 1). Assuming that such "feed-forward" loops result from authors of paper A following references of paper B, we estimate the probability $1 - d$ to follow this indirect citation path to be close to 0.5.
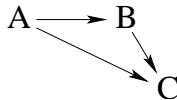

FIG. 1: Feed-forward citation loop: Publication A cites both publications B and C. About 50% of entries B in the reference list of a typical publication A cite at least one other article C in the same reference list.

To implement the Google PageRank algorithm for the citation network, we start with a uniform probability density equal to $1/N$ at each node of the network and then iterate Eq. (1). Eventually a steady state set of Google numbers for each node of the network is reached. These represent the occupation probabilities at each node for the random-walk-like process defined by Eq. (1). Finally, we sort the nodal Google numbers to determine the Google rank of each node. It is both informative and entertaining to compare the Google rank with the citation (in-degree) rank of typical and the most important publications in Physical Review.

III. GOOGLE PAGERANK FOR PHYSICAL REVIEW

Fig. 2 shows the average Google number $\langle G(k) \rangle$ for publications with k citations as a function of k. For small k, there are many publications with the same number of citations and the dispersion in $G(k)$ is small. Correspondingly, the plot of $\langle G(k) \rangle$ versus k is smooth and
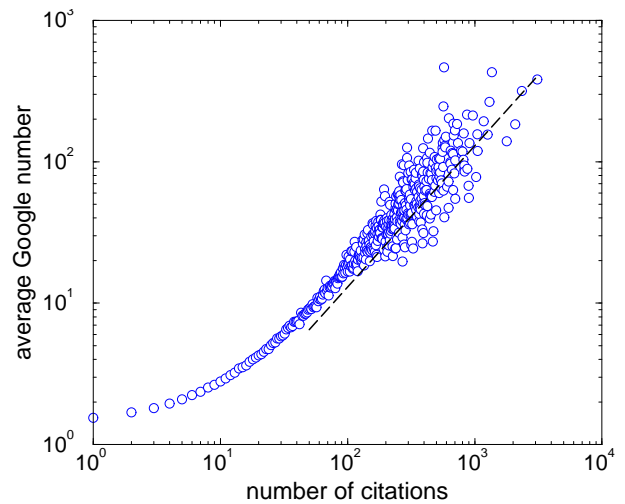

FIG. 2: Average Google number $\langle G(k) \rangle$ versus number of citations k. The dashed line of slope 1 is a guide for the eye.
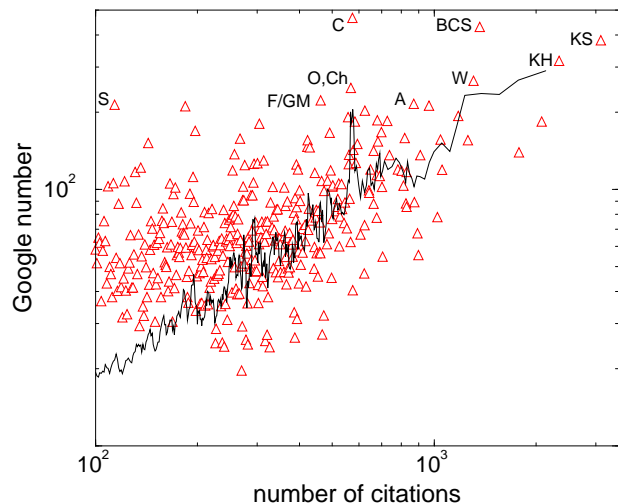

FIG. 3: Individual outlier publications. For each number of citations k, the publication with the highest Google number is plotted. The top-10 Google-ranked papers are identified by author(s) initials (see Table I). As a guide to the eye, the solid curve is a 5-point average of the data of $\langle G(k) \rangle$ versus k in Fig. 2.

increases linearly with k for $k \gtrsim 50$. Thus the average Google number and the number of citations represent similar measures of popularity, a result that has been observed previously [7, 8]. In fact, the citation and Google number distributions are qualitatively similar, further indicating that citations and Google numbers are, on the average, similar measures of importance.

However, for large k, much more interesting behavior occurs. When k is sufficiently large, there is typically only one publication with k citations. Thus instead of an average value, Fig. 3 shows the individual publications with the largest Google number for each value of cita-

tion number when $k \geq 100$. Of particular interest are the extreme outliers with respect to the linear behavior of Fig. 2. The ten articles with the highest Google numbers are shown explicitly and are identified by author initials (see Table I). Also given in Table I is the number of citations and the citation rank of these publications. While several of the highest-cited Physical Review papers appear on this list, there are also several more modestly-cited papers that are highly ranked according to the Google algorithm.

The disparity between the the Google rank and citation rank arises because, as mentioned in the previous section, the former involves both the in-degree as well as the Google PageRank of the neighboring nodes. According to the Google algorithm of Eq. (1), a citing publication ("child") j contributes a factor $\langle G_j/k_j \rangle$ to the Google number of its parent paper i. Thus for a paper to have a large Google number, its children should be important (large G_j), and also each child should have a small number of parents (small out-degree k_j). The latter ensures that the Google contribution of a child is not strongly diluted.

With this perspective, let us compare the statistical measures of the two articles "Unitary Symmetry and Leptonic Decays", Phys. Rev. Lett. **10**, 531 (1963) by N. Cabibbo (C) and "Self-Consistent Equations Including Exchange and Correlation Effects", Phys. Rev. **140**, A1133 (1965) by W. Kohn & L. J. Sham (KS). The former has the highest Google number of all Physical Review publications, while the latter is the most cited. The high Google rank of C stems from the fact that that value of $\langle G_j/k_j \rangle = 1.52 \times 10^{-6}$ for the children of C is an order of magnitude larger than the corresponding value $\langle G_j/k_j \rangle = 2.31 \times 10^{-7}$ for the children of KS. This difference more than compensates for the factor 5.6 difference in the number of citations to these two articles (3227 for KS and 574 for C as of June 2003). Looking a little deeper, the difference in $\langle G_j/k_j \rangle$ for C and KS stems from the denominator; the children of C have 15.6 citations an average, while the children of KS are slightly "better" and have 18.4 citations on average. However, the typical child of C has fewer references than a child of KS and a correspondingly larger contribution to the Google number of C.

TABLE I: The top 10 Google-ranked publications when $d = 0.5$

Google rank	Google # ($\times 10^{-4}$)	cite rank	# cites	Publication				Title	Author(s)
1	4.65	54	574	PRL	10	531	1963	Unitary Symmetry and Leptonic...	N. Cabibbo
2	4.29	5	1364	PR	108	1175	1957	Theory of Superconductivity	J. Bardeen, L. Cooper, J. Schrieffer
3	3.81	1	3227	PR	140	A1133	1965	Self-Consistent Equations...	W. Kohn & L. J. Sham
4	3.17	2	2460	PR	136	B864	1964	Inhomogeneous Electron Gas	P. Hohenberg & W. Kohn
5	2.65	6	1306	PRL	19	1264	1967	A Model of Leptons	S. Weinberg
6	2.48	55	568	PR	65	117	1944	Crystal Statistics	L. Onsager
7	2.43	56	568	RMP	15	1	1943	Stochastic Problems in...	S. Chandrasekhar
8	2.23	95	462	PR	109	193	1958	Theory of the Fermi Interaction	R. P. Feynman & M. Gell-Mann
9	2.15	17	871	PR	109	1492	1958	Absence of Diffusion in...	P. W. Anderson
10	2.13	1853	114	PR	34	1293	1929	The Theory of Complex Spectra	J. C. Slater

The remaining research articles on the top-10 Google-rank list but outside the top-10 citation list are easily recognizable as seminal publications. For example, Onsager's 1944 paper presents the exact solution of the two-dimensional Ising model; both a calculational *tour de force*, as well as a central development in the theory of critical phenomena. The paper by Feynman and Gell-Mann introduced the $V - A$ theory of weak interactions that incorporated parity non-conservation and became the "standard model" of weak interactions. Anderson's paper, "Absence of Diffusion in Certain Random Lattices" gave birth to the field of localization and is cited by the Nobel prize committee for the 1977 Nobel prize in physics.

The last entry in the top-10 Google-rank list, "The Theory of Complex Spectra", by J. C. Slater (S) is par-

ticularly striking. This article has relatively few citations (114 as of June 2003) and a relatively low citation rank (1853$^{\text{th}}$), but its Google number 2.13×10^{-4} is only a factor 2.2 smaller than that of Cabibbo's paper! What accounts for this high Google rank? From the scientific standpoint, Slater's paper introduced the determinant form for the many-body wavefunction. This form is so ubiquitous in current literature that very few articles actually cite the original work when the Slater determinant is used. The Google PageRank algorithm identifies this hidden gem primarily because the average Google contribution of the children of S is $\langle G_j/k_j \rangle = 3.51 \times 10^{-6}$, which is a factor 2.3 larger than the contribution of the children of C. That is, the children of Slater's paper were both influential and Slater loomed as a very important father figure to his children.

TABLE II: The remaining top-100 Google-ranked papers when $d = 0.5$
in which the ratio of Google rank to citation rank is greater than 10.

Google rank	Google # ($\times 10^{-4}$)	cite rank	# cites	Publication			Title	Author(s)
1	4.65	54	574	PRL	10	531	1963 Unitary Symmetry and Leptonic...	N. Cabibbo
8	2.23	95	462	PR	109	193	1958 Theory of the Fermi Interaction	R. P. Feynman & M. Gell-Mann
10	2.13	1853	114	PR	34	1293	1929 The Theory of Complex Spectra	J. C. Slater
12	2.11	712	186	PR0	43	804	1933 On the Constitution of...	E. Wigner & F. Seitz
20	1.80	228	308	PR0	106	364	1957 Correlation Energy of an ...	M. Gell-Mann & K. Brueckner
21	1.69	616	198	PRL	58	408	1987 Bulk superconductivity at ...	R. J. Cava et al.
25	1.58	311	271	PRL	58	405	1987 Evidence for superconductivity ...	C. W. Chu et al.
30	1.51	1193	144	PRL	10	84	1963 Photon Correlations	R. J. Glauber
35	1.42	12897	39	PR0	35	509	1930 Cohesion in Monovalent Metals	J. C. Slater
49	1.21	1342	136	PR0	60	252	1941 Statistics of the Two- ...	H. A. Kramers & G. H. Wannier
58	1.17	1433	135	PR0	81	440	1951 Interaction Between the ...	C. Zener
59	1.17	5196	66	PR0	45	794	1934 Electronic Energy Bands in ...	J. C. Slater
60	1.16	2927	108	PRB	28	4227	1983 Electronic structure of ...	L. F. Mattheiss & D. R. Hamann
64	1.12	642	199	PR0	52	191	1937 The Structure of Electronic ...	G. H. Wannier
70	1.08	1653	130	PRL	10	518	1963 Classification of Two-Electron ...	J. Cooper, U. Fano & F. Prats
72	1.06	1901	118	PR0	46	509	1934 On the Constitution of ...	E. Wigner & F. Seitz
73	1.05	876	180	PR0	75	486	1949 The Radiation Theories of ...	F. J. Dyson
78	1.03	1995	119	PR0	109	1860	1958 Chirality Invariance and ...	E. Sudarshan & R. Marshak
85	1.00	201853	3	PRB	22	5797	1980 Cluster formation in ...	H. Rosenstock & C. Marquardt
87	0.99	10168	48	PRL	6	106	1961 Population Inversion and ...	A. Javan, W. Bennett, D. Herriott
90	0.98	3231	86	PR0	79	350	1950 Antiferromagnetism. ...	P. W. Anderson
92	0.97	1199	149	PR0	76	749	1949 The Theory of Positrons	R. P. Feynman

The striking ability of the Google PageRank algorithm to identify influential papers can be seen when we consider the top-100 Google-ranked papers. Table II shows the subset of publications on the top-100 Google rank in which the ratio of Google rank to citation rank is greater than 10; that is, publications with anomalously high Google rank compared to their citation rank. This list contains many easily-recognizable papers for the average physicist. For example, the publication by Wigner and Seitz, "On the Constitution of Metallic Sodium" introduced Wigner-Seitz cells, a construction that appears in any solid-state physics text. The paper by Gell-Mann and Brueckner, "Correlation Energy of an Electron Gas at High Density" is a seminal publication in many-body theory. The publication by Glauber, "Photon Correlations", was recognized for the 2005 Nobel prize in physics. The Kramers-Wannier article, "Statistics of the Two-Dimensional Ferromagnet. Part I", showed that a phase transition occurs in two dimensions, contradicting the common belief at the time. The article by Dyson, "The Radiation Theories of Tomonaga, Schwinger, and Feynman", unified the leading formalisms for quantum electrodynamics and it is plausible that this publication would have earned Dyson the Nobel prize if it could have been shared among four individuals. One can offer similar rationalizations for the remaining articles in this table.

On the other hand, an apparent mistake is the paper by Rosenstock and Marquardt, "Cluster formation in two-dimensional random walks: Application to photolysis of silver halides" (RM). Notice that this article has only 3 citations! Why does RM appear among the top-100 Google-ranked publications? In RM, a model that is essentially diffusion-limited aggregation is introduced. Although these authors had stumbled upon a now-famous model, they focused on the kinetics of the system and apparently did not appreciate its wonderful geometrical features. This discovery was left to one of the children of RM—the famous paper by T. Witten and L. Sander, "Diffusion-Limited Aggregation, a Kinetic Critical Phenomenon" Phys. Rev. Lett. **47**, 1400 (1981), with 680 citations as of June 2003. Furthermore, the Witten and Sander article has only 10 references; thus a substantial fraction of its fame is exported to RM by the Google PageRank algorithm. The appearance of RM on the list of top-100 Google-ranked papers occurs precisely because of the mechanics of the Google PageRank algorithm in which being one of the few references of a famous paper makes a huge contribution to the Google number.

A natural question to ask is whether the Google rankings are robust with respect to the value of the free parameter d in the Google algorithm. As mentioned above, we believe that our *ad hoc* choice of $d = 0.5$ accounts in a reasonable way for the manner in which citations are actually made. For $d = 0.15$, as in the original Google algorithm, the Google rankings of highly-cited papers locally reorder to a considerable extent compared to the rankings for the case $d = 0.5$, but there is little global reordering. For example, all of the top 10 Google-ranked papers calculated with $d = 0.5$ remained among the top-
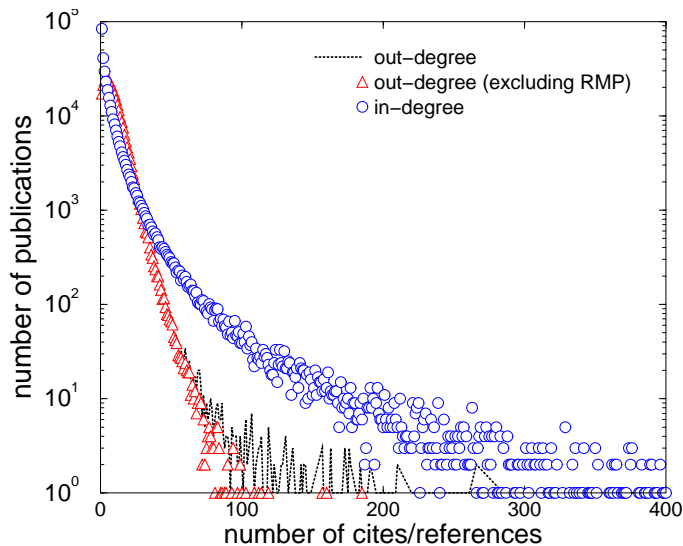
FIG. 4: The in-degree distribution (citations to) and out-degree distribution (references from) for all Physical Review publications. The out-degree distribution is shown with and without the contribution of Reviews of Modern Physics.

50 Google-ranked papers for $d = 0.15$. Thus up to this degree of variation, Google rankings are a robust measure.

On the other hand, in the limit $d \to 1$ Google rank approaches citation rank. For example, for $d = 0.9$, 7 of the top-10 Google-ranked papers with $d = 0.9$ are also among the 10 most-cited particles, while the citation ranks of the remaining 3 of the top-10 Google-ranked articles are 19, 54, and 56. In fact, we argue that the Google rank reduces to citation rank as $d \to 1$. To show this, we first notice that in the extreme case of $d = 1$, the Google number of each node equals $1/N$. For $d \to 1$, we therefore write $d = 1 - \epsilon$, with $\epsilon \ll 1$, and also assume that there is a correspondingly small deviation of the Google numbers from $1/N$. Thus we write $G_i = \frac{1}{N} + \mathcal{O}(\epsilon)$. Substituting these into Eq. (1), we obtain

$$
\begin{aligned}
G_i &= \epsilon \sum_j \frac{G_j}{k_j} + \frac{1 - \epsilon}{N} \\
&\approx \frac{1}{N} \left[1 + \epsilon \left(\sum_j \frac{1}{k_j} - 1 \right) \right]
\end{aligned}
\tag{2}
$$

To estimate the sum in Eq. (2), we use the fact that the out-degree distribution is relatively narrow (Fig. 4), especially if we exclude the broad tail that is caused by the contributions of review articles that appear in the Reviews of Modern Physics. While the mean in-degree and out-degrees are both close to 9 (and should be exactly equal for the complete citation network), the dispersion for the in degree is 23.15, while the dispersion for the out degree (excluding Reviews of Modern Physics) is 8.64.

As a result of the sharpness of the out-degree distribu-

tion, the sum $\sum_{j \text{ nn } i} \frac{1}{k_j}$ for nodes with high in-degree is approximately equal to the in-degree d_i of node i times $\langle \frac{1}{k} \rangle$. With this assumption, Eq. (2) becomes

$$
G_i = \frac{1}{N} \left[1 + \epsilon \left(d_i \left\langle \frac{1}{k} \right\rangle - 1 \right) \right] .
\tag{3}
$$

That is, the leading correction to the limiting $d = 1$ result that $G_i = \frac{1}{N}$ is proportional to the in-degree of each node. Thus as $d \to 1$, the Google rank of each node is identical to its citation rank under the approximation that we neglect the effect of the dispersion of the out-degree in the citation network.

IV. CONCLUSIONS

We believe that protocols based on the Google PageRank algorithm hold a great promise for quantifying the impact of scientific publications. They provide a meaningful extension to traditionally-used importance measures, such as the number of citation of individual articles and the impact factor for journals as a whole. The PageRank algorithm implements, in an extremely simple way, the reasonable notion that citations from more important publications should contribute more to the rank of the cited paper than those from less important ones. Other ways of attributing a quality for a citation would require much more detailed contextual information about the citation itself.

The situation in citation networks is not that dissimilar from that in the World Wide Web, where hyperlinks contained in popular websites and pointing to your webpage would bring more Internet traffic to you and thus would contribute substantially to the popularity of your own webpage. Scientists commonly discover relevant publications by simply following chains of citation links from other papers. Thus it is reasonable to assume that the popularity or "citability" of papers may be well approximated by the random surfer model that underlies the PageRank algorithm. One meaningful difference between the WWW and citation networks is that citation links cannot be updated after publication, while WWW hyperlinks keep evolving together with the webpage containing them. Thus scientific papers and their citations tend to age much more rapidly than active webpages. These differences could be taking into account by explicitly incorporating the effects of aging into the Page Rank algorithm [10].

[1] D. J. de Solla Price, Science **149**, 510 (1965); J. Amer. Soc. Inform. Sci. **27**, 292 (1976).

[2] J. Laherrère and D. Sornette, Eur. Phys. J. B **2**, 525 (1998).

[3] S. Redner, Eur. Phys. J. B **4**, 131 (1998).

[4] S. Redner, *Physics Today* **58**, 49 (2005); see http://arxiv.org/abs/physics/0407137 for more detailed information.

[5] S. Brin and L. Page, Computer Networks and ISDN Systems, **30**, 107 (1998).

[6] See *e.g.*, J. Bollen, M. A. Rodriguez, H. Van de Sompel cs.DL/0601030, which was later put in the context of other relevant publications in P. Ball, Nature **439**, 770–771 (2006).

[7] S. Fortunato, M. Boguna, A. Flammini, and F. Menczer, cs.IR/0511016.

[8] S. Fortunato, A. Flammini, and F. Menczer, cond-mat/0602081.

[9] The actual fraction of "followed citations" (such as B in Fig. 1) is 42% for the entire dataset and 51% for papers published during the last 4 years.

[10] H. Xie, D. Walker, K.-K Yan, P. Chen, S. Redner, S. Maslov, in preparation.